



# Distributed Systems

## Consistency & Replication I

Olaf Landsiedel

# Last Time

- Logical Clocks
- Vector Clocks
- Consistent Snapshots
- Passive Monitoring
- Active Monitoring

# Active Monitoring

- Recall: we used vector clocks so we can construct consistent global states (or cuts) from processes' local histories via passive monitoring
- Remember, there is another type of monitoring – **active monitoring**
  - Active monitoring is also referred to as constructing a ***distributed snapshot***

# Distributed Snapshot Protocol (Chandy and Lamport)

- Monitor  $p_0$  sends “take snapshot” message to all processes  $p_i$
- If  $p_i$  receives such “take snapshot” message for the first time, it:
  - Records its state
  - Stops doing any activity related to the distributed computation
  - Relays the “take snapshot” message on all of its outgoing channels
  - Starts recording a state of its incoming channels – records all messages that have been delivered after the receipt of the first “take snapshot” message

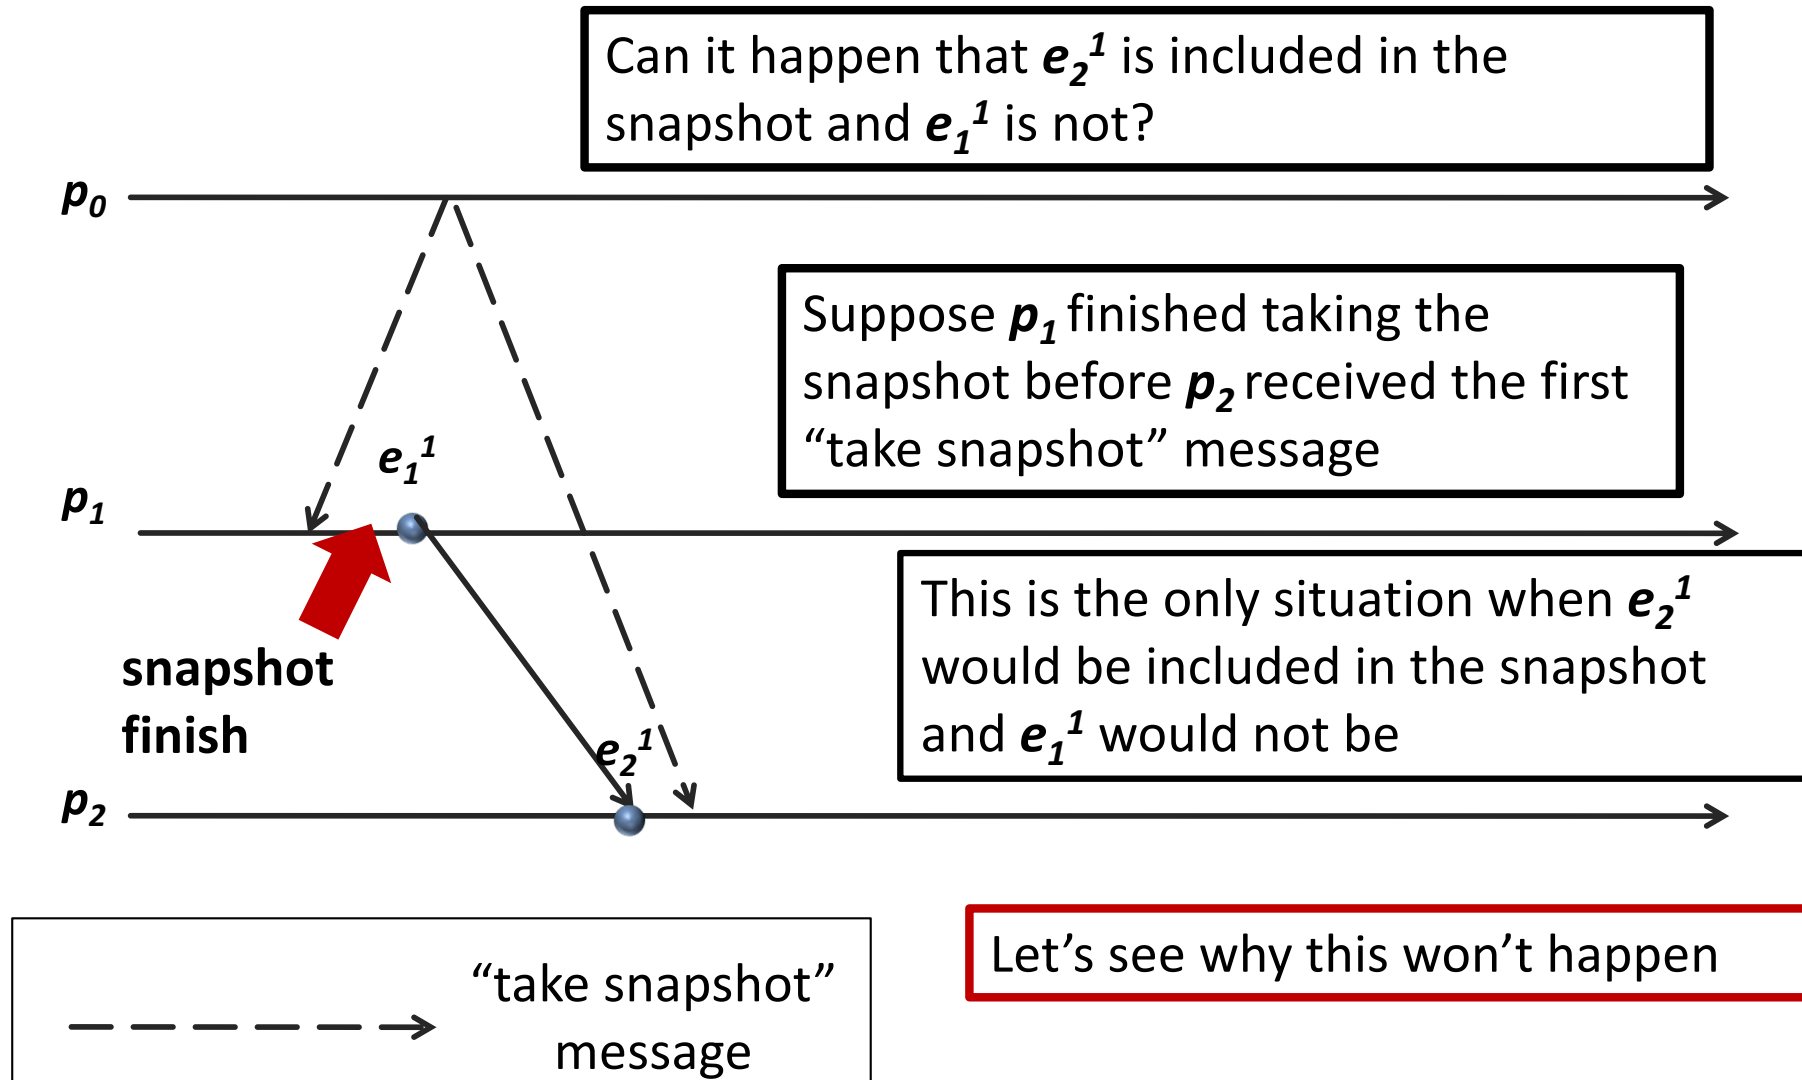
# Distributed Snapshot Protocol (Chandy and Lamport)(cont.)

- When  $p_i$  receives a “take snapshot” message from process  $p_j$ , it stops recording the state of the channel between itself and  $p_j$
- When  $p_i$  received “take snapshot” message from all processes and from  $p_o$ , it stops recording the snapshot and sends it to  $p_o$

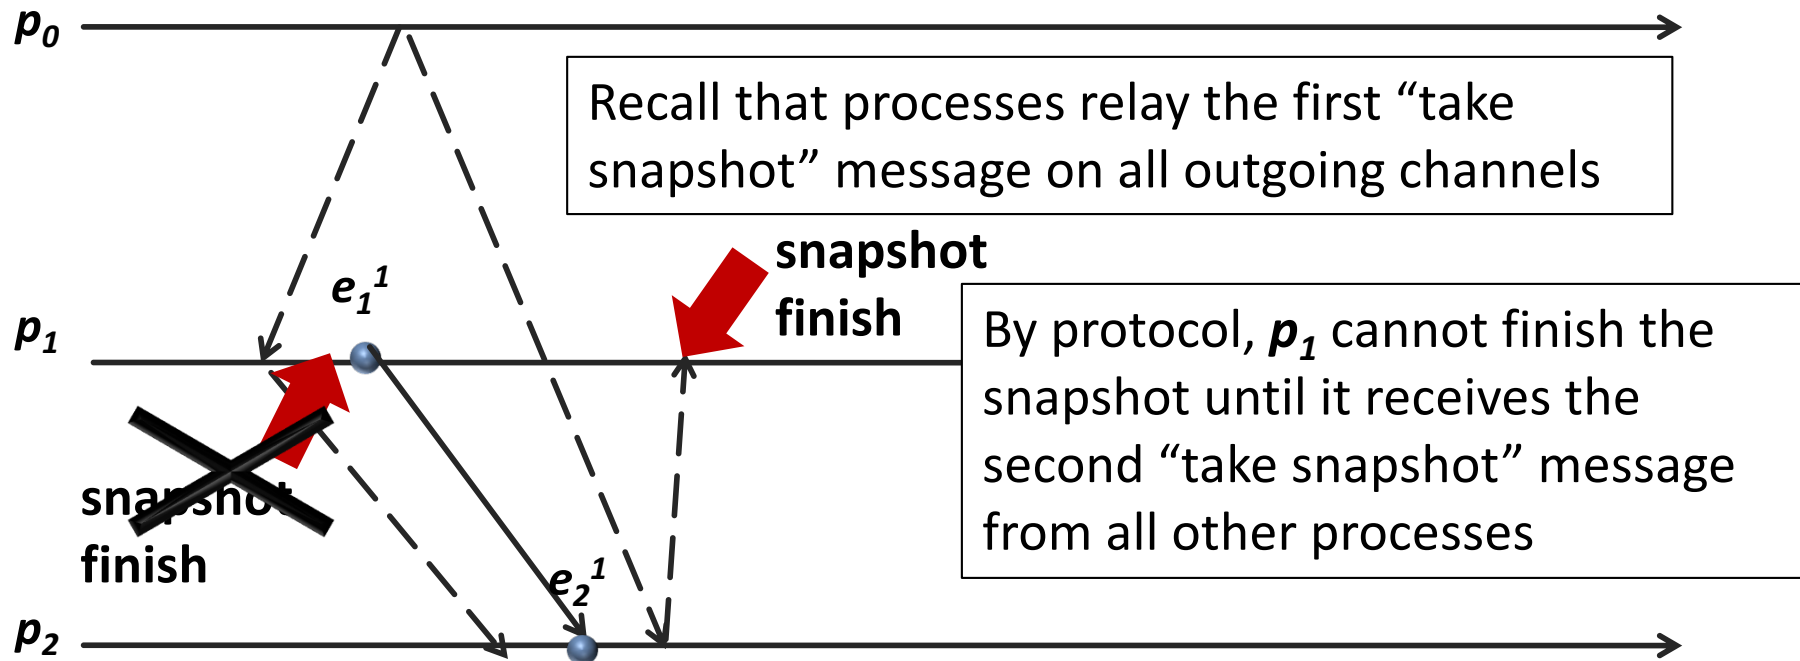
# Properties of Chandy-Lamport Protocol

- Let's see why this protocol constructs only consistent snapshots, provided that FIFO channels are used
- Recall the key property of a consistent snapshot (i.e., consistent cut):
  - If event  $(e \rightarrow e'')$  and  $e''$  is included in the snapshot, then  $e$  must also be included in that snapshot
- Let's see why this property holds if we use the Chandy-Lamport distributed snapshot protocol

# Properties of Chandy-Lamport Protocol



# Chandy-Lamport Protocol: Example 1




So  $p_1$  cannot send  $e_1^1$  and finish taking the snapshot before  $p_2$  received the first "take snapshot" message.

So that situation will not happen



# Chandy-Lamport Protocol: Example 2



But  $p_2$  must **stop recording events** from  $p_1$  after it receives the "take snapshot" message from  $p_1$ .

# Questions

- Active monitoring vs. passive monitoring?
  - Active monitoring: query the system
- Goal of Chandy and Lamport protocol?
  - Global snapshot

# Summary

- Constructing a consistent global state is difficult in absence of global clocks
- Consistent global states can be constructed using
  - Active monitoring (distributed snapshots)
  - Passive monitoring (local event histories with vector clock timestamps)

# Today...

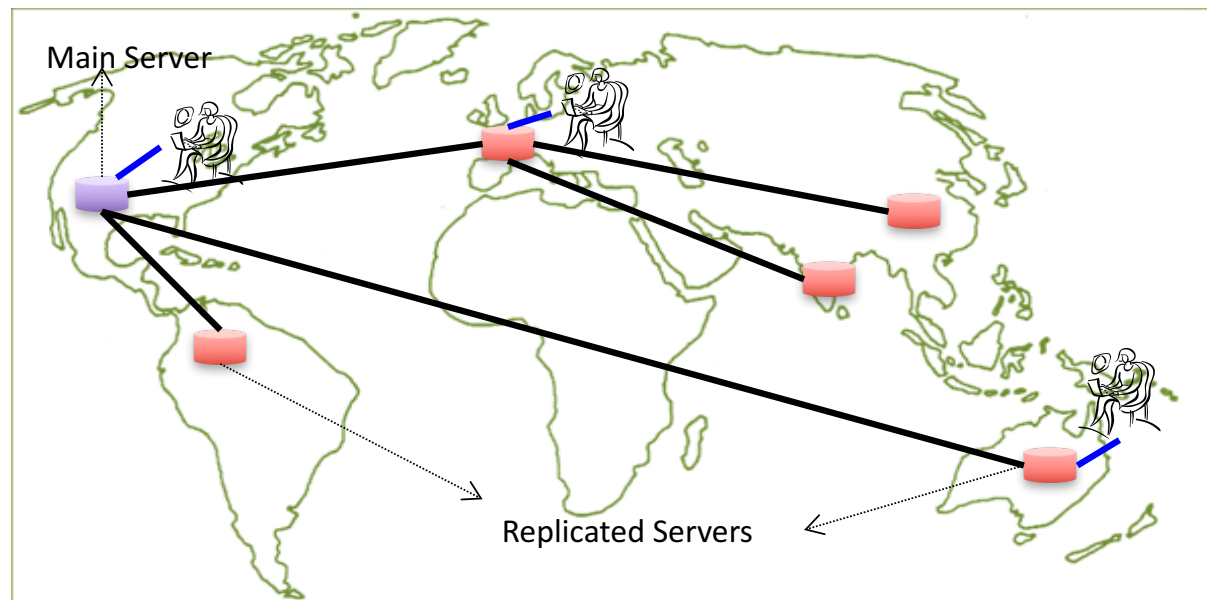
- Consistency and Replication
  - Introduction
  - Data-centric Consistency Models

# Why Replication?

- Replication is the process of maintaining the data at multiple computers
- Replication is necessary for:
  - **Improving performance**
    - A client can access the replicated copy of the data that is near to its location
  - **Increasing the availability of services**
    - Replication can mask failures such as server crashes and network disconnection
  - **Enhancing the scalability of the system**
    - Requests to the data can be distributed to many servers which contain replicated copies of the data
  - **Securing against malicious attacks**
    - Even if some replicas are malicious, secure data can be guaranteed to the client by relying on the replicated copies at the non-compromised servers


# 1. Replication for Improving Performance

- Example Applications
  - Caching webpages at the client browser
  - Caching IP addresses at clients and DNS Name Servers
  - Caching in Content Delivery Network (CDNs)
    - Commonly accessed contents, such as software and streaming media, are cached at various network locations



## 2. Replication for High-Availability

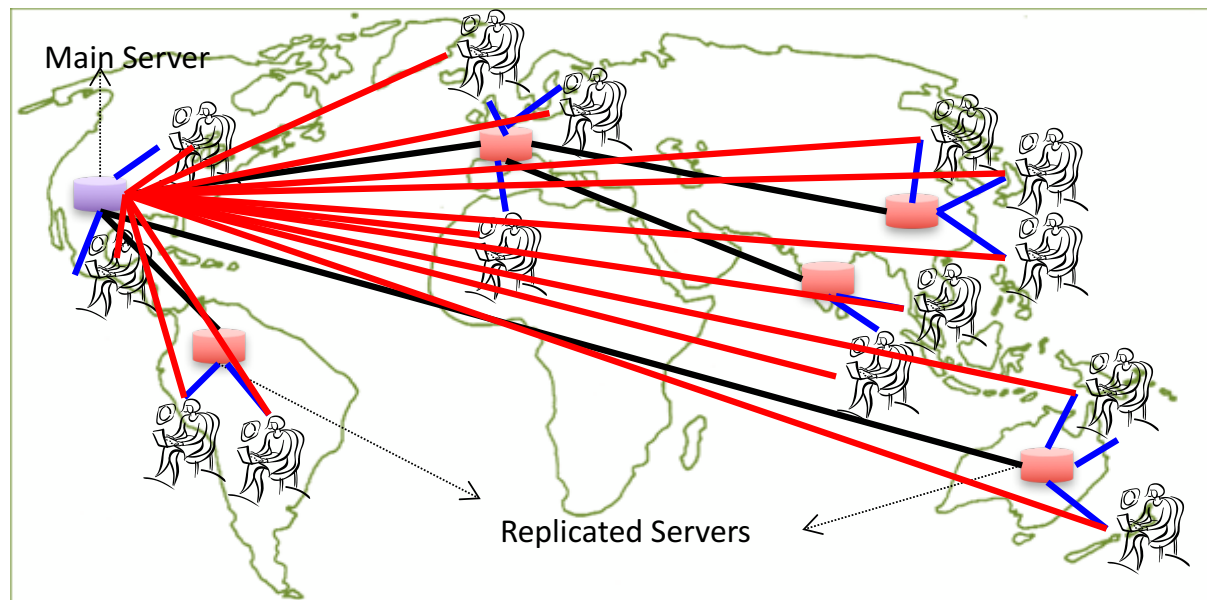
- Availability can be increased by storing the data at replicated locations (instead of storing one copy of the data at a server)



- Example:  
Google File-System and Chubby replicate the data at computers across different racks, clusters and data-centers
  - If one computer or a rack or a cluster crashes, then a replica of the data can still be accessed

### 3. Replication for Enhancing Scalability


- Distributing the data across replicated servers helps in avoiding bottle-necks at the main server
  - It balances the load between the main and the replicated servers
- Example: Content Delivery Networks decrease the load on main servers of the website





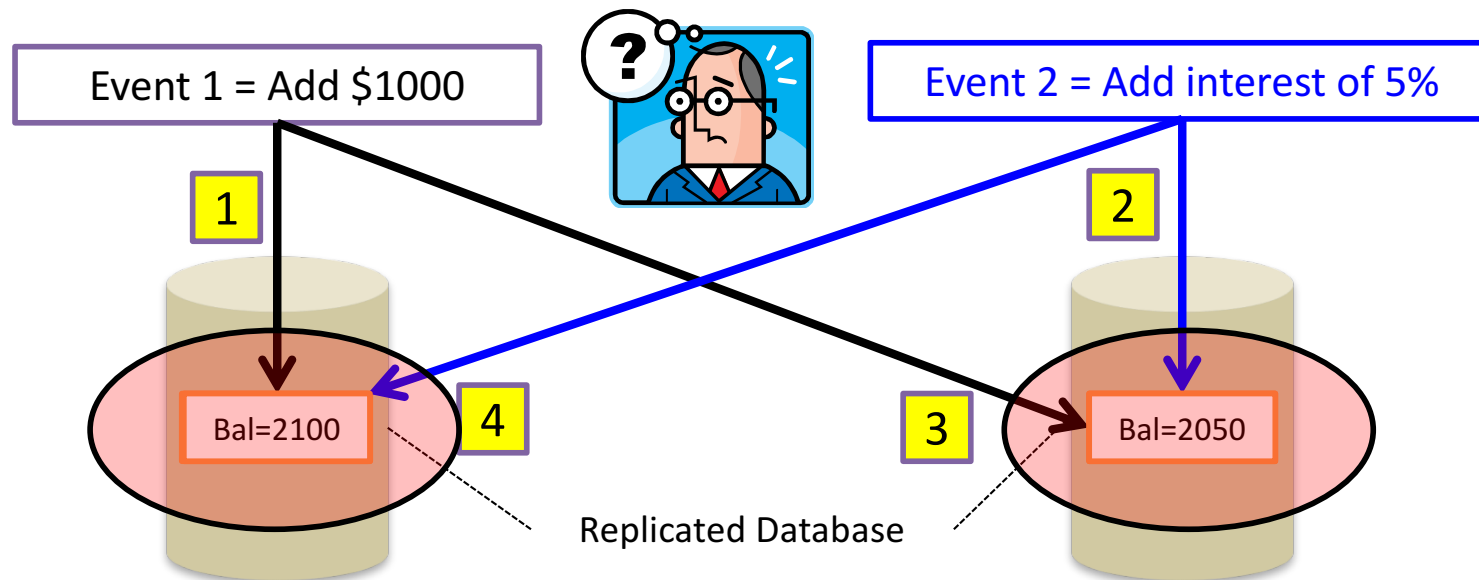
## 4. Replication for Securing Against Malicious Attacks

- If a minority of the servers that hold the data are malicious, the non-malicious servers can outvote the malicious servers, thus providing security.
- The technique can also be used to provide fault-tolerance against non-malicious but faulty servers



# Why Consistency?

- In a DS with replicated data, one of the main problems is keeping the data consistent
- An example:
  - In an e-commerce application, the bank database has been replicated across two servers
  - Maintaining consistency of replicated data is a challenge



# Overview of Consistency and Replication

## Today's lecture

- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
  - When, where and by whom replicas should be placed?
  - Which consistency model to use for keeping replicas consistent?
- Consistency Protocols
  - We study various implementations of consistency models


## Next lectures

# Overview

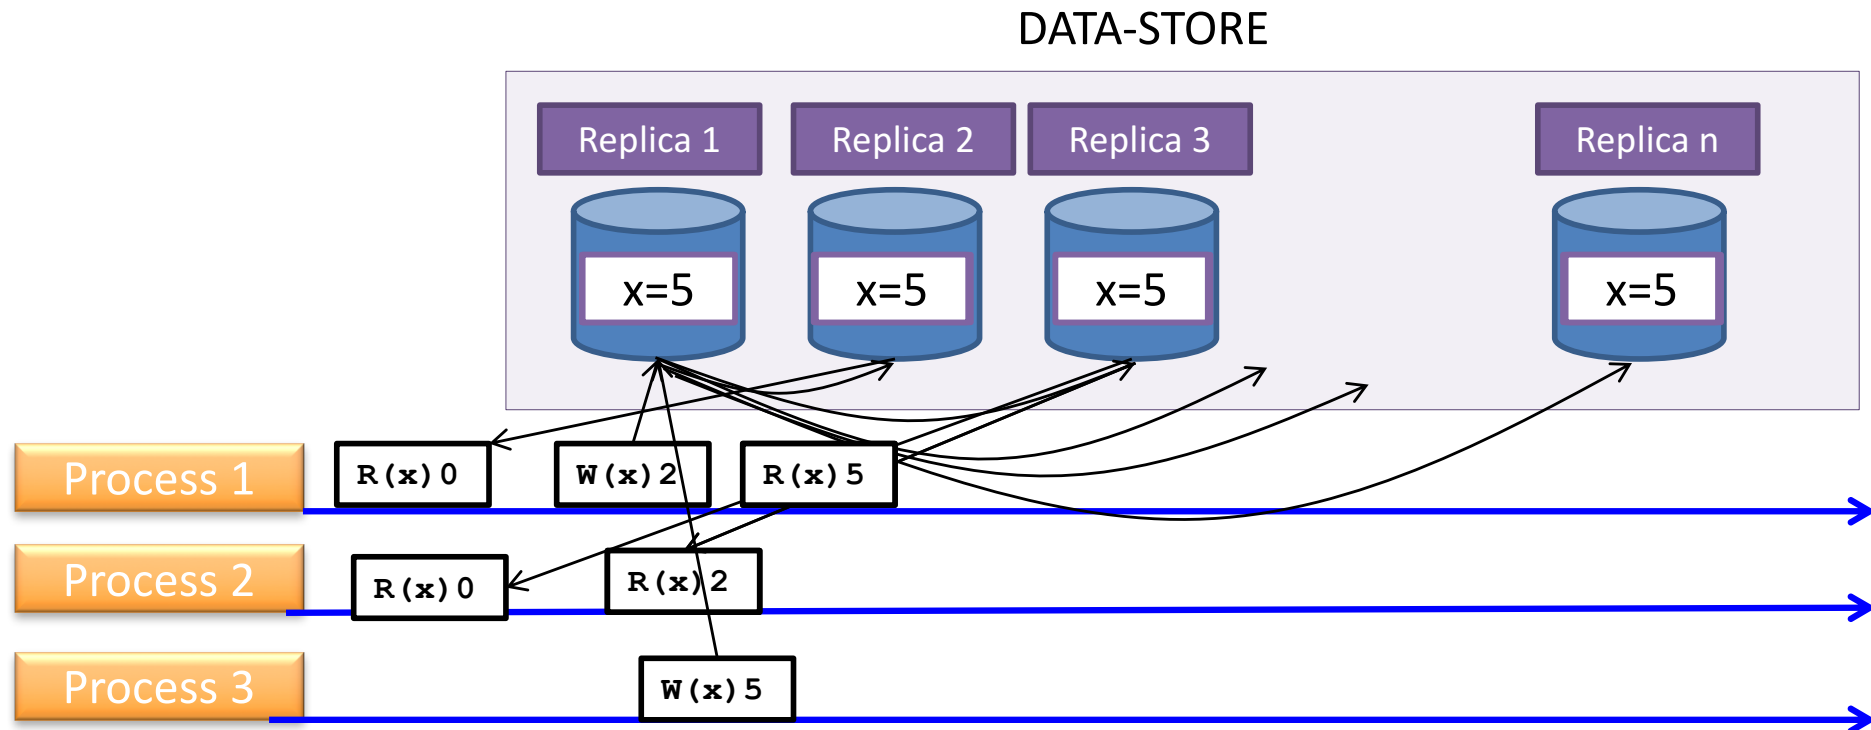
- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

# Introduction to Consistency and Replication

- In a distributed system, shared data is typically stored in distributed shared memory, distributed databases or distributed file systems.
  - The storage can be distributed across multiple computers
  - Simply, we refer to a series of such data storage units as data-stores
- Multiple processes can access shared data by accessing any replica on the data-store
  - Processes generally perform read and write operations on the replicas



# Maintaining Consistency of Replicated Data



## Strict Consistency

- Data is always fresh
  - After a write operation, the update is propagated to all the replicas
  - A read operation will result in reading the most recent write
- If there are many writes, this leads to large overheads

**P1** = Process P1

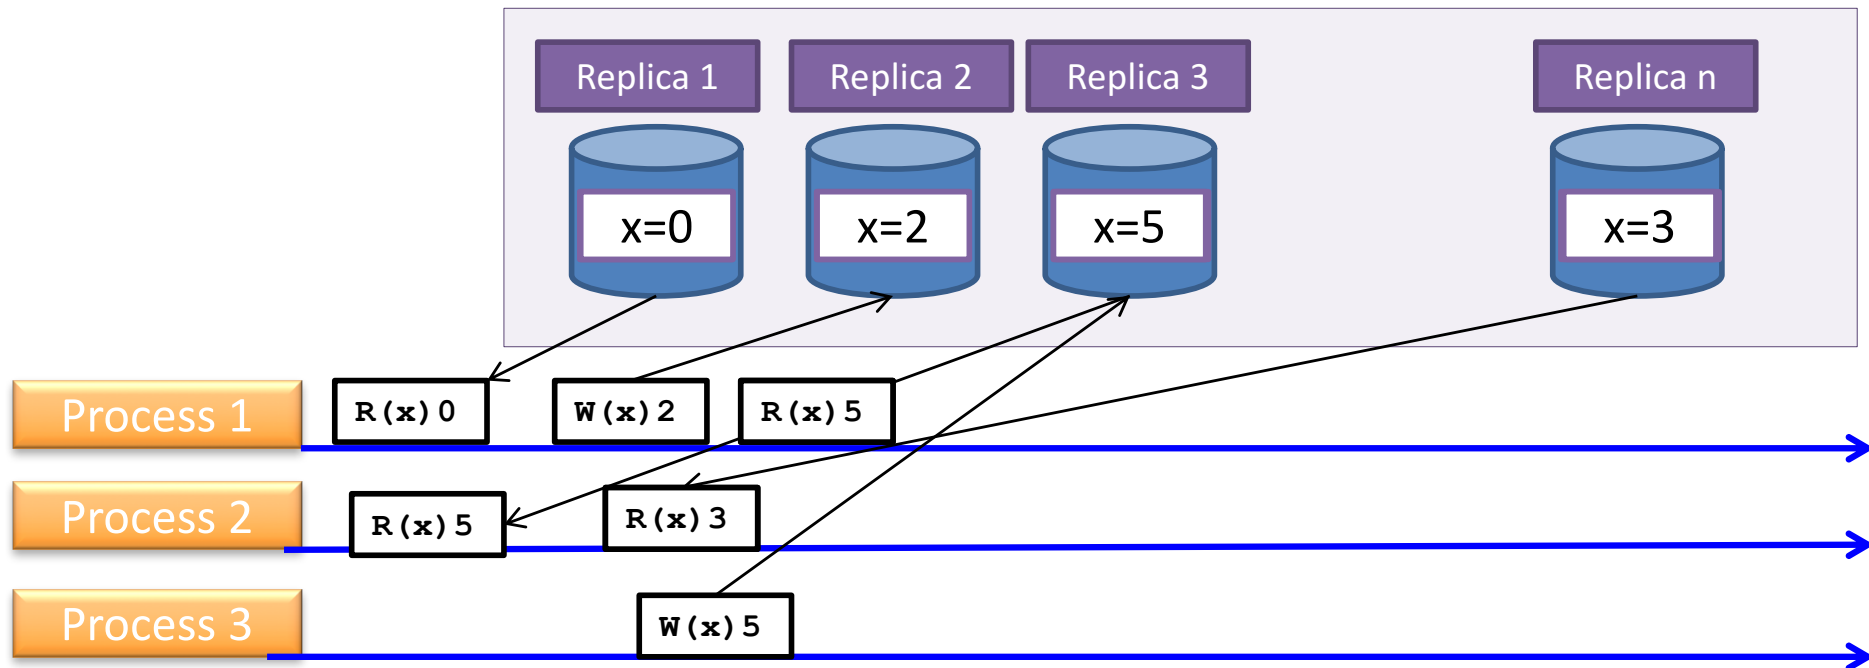
$\longrightarrow$  = Timeline at P1

**$R(x) b$**  = Read variable  $x$ ;  
Result is  $b$

**$W(x) b$**  = Write variable  $x$ ;  
Result is  $b$

# Maintaining Consistency of Replicated Data (cont'd)

DATA-STORE



## Loose Consistency

- Data might be stale
  - A read operation may result in reading a value that was written long back
  - Replicas are generally out-of-sync
- The replicas may sync at coarse grained time, thus reducing the overhead

**P1** = Process P1


**→** = Timeline at P1

**R(x) b** = Read variable x;  
Result is b

**W(x) b** = Write variable x;  
Result is b

# Trade-offs in Maintaining Consistency

- Maintaining consistency should balance between the strictness of consistency versus efficiency
  - How much consistency to we need?
  - Good-enough consistency depends on your application





# Questions

- Strict consistency vs. loose consistency?
  - Strict consistency: data is always up-to-date
  - Loose consistency: data maybe be outdated
- Which one to implement for
  - a Bank?
  - Facebook?

# Consistency Model

- A consistency model is a contract between
  - the process that wants to use the data, and
  - the replicated data repository (or data-store)
- A consistency model states the level of consistency provided by the data-store to the processes while reading and writing the data

# Types of Consistency Models

- Consistency models can be divided into two types:
  - Data-Centric Consistency Models
    - These models define how the data updates are propagated across the replicas to keep them consistent
  - Client-Centric Consistency Models
    - These models assume that clients connect to different replicas at each time
    - The models ensure that whenever a client connects to a replica, the replica is brought up to date with the replica that the client accessed previously

# Overview

- Consistency Models
  - Data-Centric Consistency Models
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

# Data-Centric Consistency Models

- Data-centric Consistency Models describe how the replicated data is kept consistent, and what the process can expect
- Under Data-centric Consistency Models, we study two types of models:
  - Consistency Specification Models:
    - These models enable specifying the consistency levels that are tolerable to the application
  - Models for Consistent Ordering of Operations:
    - These models specify the order in which the data updates are propagated to different replicas

# Overview

- Consistency Models
  - Data-Centric Consistency Models
    - Consistency Specification Models
    - Models for Consistent Ordering of Operations
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

# Consistency Specification Models

- In replicated data-stores, there should be a mechanism to:
  - Measure how inconsistent the data might be on different replicas
  - How replicas and applications can specify the tolerable inconsistency levels
- Consistency Specification Models enable measuring and specifying the level of inconsistency in a replicated data-store
- We study a Consistency Specification Model called Continuous Consistency Model

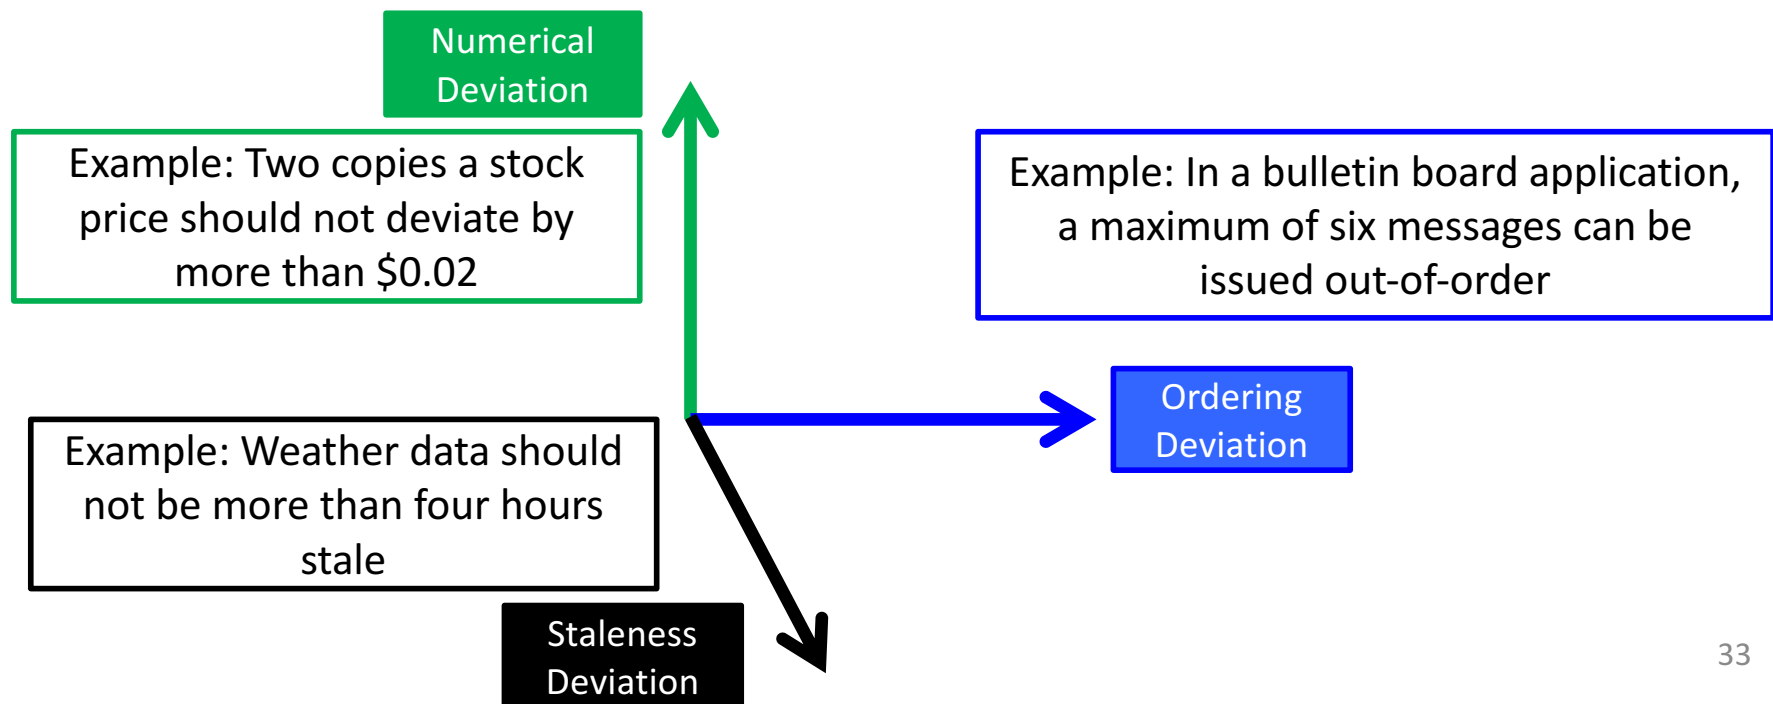
# Continuous Consistency Model

- Continuous Consistency Model
  - is used to measure inconsistencies and
  - express what inconsistencies can be expected in the system



# Continuous Consistency Ranges

- Level of consistency is defined over three independent axes:
  - **Numerical Deviation**: Deviation in the numerical values between replicas
  - **Order Deviation**: Deviation with respect to the ordering of update operations
  - **Staleness Deviation**: Deviation in the staleness between replicas



# Consistency Unit (Conit)

- Consistency unit (Conit) specifies the data unit over which consistency is measured
  - For example, conit can be defined as a record representing a single stock
- Level of consistency is measured by each replica along the three dimensions
  - Numerical Deviation
    - For a given replica R, how many updates at other replicas are not yet seen at R? What is the effect of the non-propagated updates on local Conit values?
  - Order Deviation
    - For a given replica R, how many local updates are not propagated to other replicas?
  - Staleness Deviation
    - For a given replica R, how long has it been since updates were propagated?

# Conit and Consistent

Assume

Order De  
not prese

There is one operation at B which is not yet pushed to A with a value difference of  $x=2, y=0$

There are no operations at A which are not yet pushed to B

No operations at A are not yet pushed to B

One operation at B was not yet pushed to A

Replica A

$x; y$

Operation

Result

<5,B>	$x+=2$	$x=2$
<10,A>	$y+=1$	$y=1$
<14,A>	$x+=1$	$x=3$
<23,A>	$y+=3$	$y=4$

Replica B

$x; y$

Operation

Result

<5,B>	$x+=2$	$x=2$
<16,B>	$y+=1$	$y=1$

Replica A					Replica B				
x	y	VC	Ord	n(w)	x	y	VC	Ord	n(w)
0	0	(0,0)	0	0(0,0)	0	0	(0,0)	0	0(0,0)
0	0	(0,0)	0	1(2,0)	2	0	(0,5)	1	0(0,0)
2	0	(1,5)	0	0(0,0)	2	0	(0,5)	0	0(0,0)
2	1	(10,5)	1	0(0,0)	2	0	(0,5)	0	1(0,1)
2	1	(10,5)	1	1(0,1)	2	1	(0,16)	1	1(0,1)
3	1	(14,5)	2	1(0,1)	2	1	(0,16)	1	2(1,1)
3	4	(23,5)	3	1(0,1)	2	1	(0,16)	1	3(1,4)

<5,B> =

Operation performed at B when the vector clock was 5

<m,n>

= Uncommitted operation

<m,n>

= Committed operation

$x; y$

= A Conit

# FOR PRINTING – WITHOUT BUBBLES

Assume a **global observer** keeps a **table of operations** and measures

Order Deviation at a replica R is the number of operations in R that are not present at the other replicas

Numerical Deviation at replica R is defined as  $n(w)$ , where

$n$  = # of operations at other replicas that are not yet seen by R,

$w$  = weight of the deviation

= list of update amount for all variables in a Conit

Replica A					Replica B				
x	y	VC	Ord	$n(w)$	x	y	VC	Ord	$n(w)$
0	0	(0,0)	0	0(0,0)	0	0	(0,0)	0	0(0,0)
0	0	(0,0)	0	1(2,0)	2	0	(0,5)	1	0(0,0)
2	0	(1,5)	0	0(0,0)	2	0	(0,5)	0	0(0,0)
2	1	(10,5)	1	0(0,0)	2	0	(0,5)	0	1(0,1)
2	1	(10,5)	1	1(0,1)	2	1	(0,16)	1	1(0,1)
3	1	(14,5)	2	1(0,1)	2	1	(0,16)	1	2(1,1)
3	4	(23,5)	3	1(0,1)	2	1	(0,16)	1	3(1,4)

Replica A

x; y

Operation

Result

<5,B>

x+=2

x=2

<10,A>

y+=1

y=1

<14,A>

x+=1

x=3

<23,A>

y+=3

y=4

Replica B

x; y

Operation

Result

<5,B>

x+=2

x=2

<16,B>

y+=1

y=1

<5,B> =

Operation performed at B  
when the vector clock was 5

<m,n>

= Uncommitted  
operation

<m,n>

= Committed  
operation

x;y

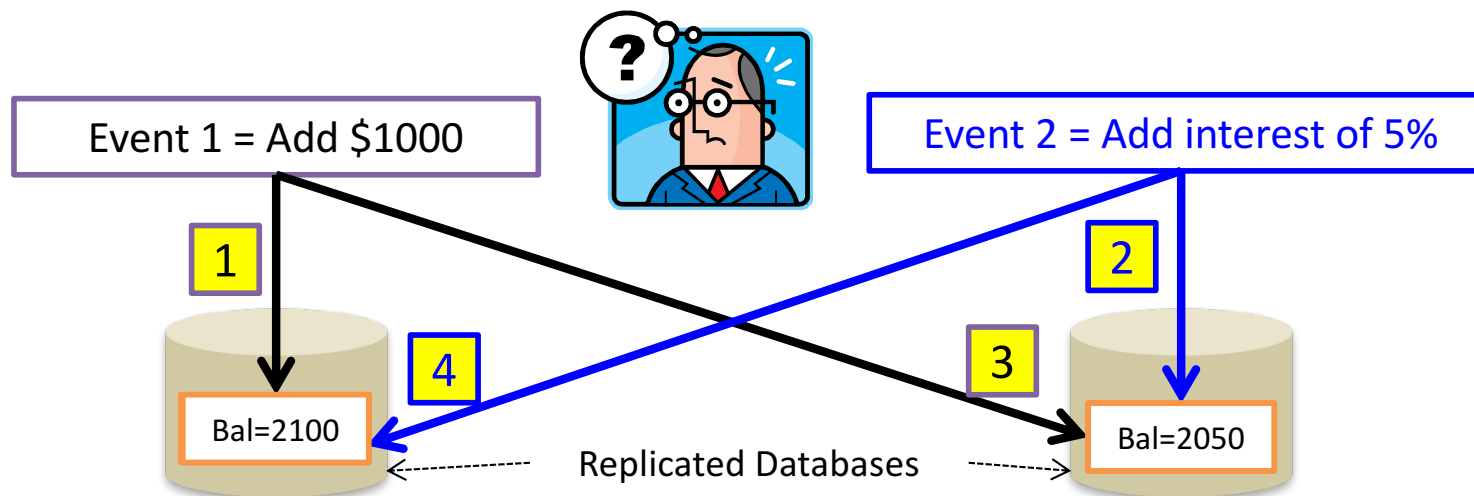
= A Conit

# Overview

- Consistency Models
  - Data-Centric Consistency Models
    - Continuous Specification Models
    - Models for Consistent Ordering of Operations
  - Client-Centric Consistency Models
- Replica Management
- Consistency Protocols

# Why is Consistent Ordering Required in Replication?

- In several applications, the order or the sequence in which the replicas commit to the data store is critical
- Example:



- Continuous Specification Models defined how inconsistency is measured
  - However, the models did not enforce any order in which the data is committed

# Consistent Ordering of Operations (cont'd)

- Whenever a replica is updated, it propagates the updates to other replicas at some point in time
- Updating different replicas is carried out by passing messages between the replica data-stores
- We will study different types of ordering and consistency models arising from these orderings

# Types of Ordering

- We will study three types of ordering of messages that meet the needs of different applications:
  - Total Ordering
  - Sequential Ordering
    - Sequential Consistency Model
  - Causal Ordering
    - Causal Consistency Model



# Questions

- How do we measure consistency?
  - Global observer
- Metrics to measure consistency?
  - **Numerical Deviation:** Deviation in the numerical values between replicas
  - **Order Deviation:** Deviation with respect to the ordering of update operations
  - **Staleness Deviation:** Deviation in the staleness between replicas

# Types of Ordering

- Total Ordering
- Sequential Ordering
- Causal Ordering



# Total Ordering

## – Total Order

- If process  $P_i$  sends a message  $m_i$  and  $P_j$  sends  $m_j$ , and if one correct process delivers  $m_i$  before  $m_j$  then every correct process delivers  $m_i$  before  $m_j$
- Messages can contain replica updates, such as passing the read or write operation that needs to be performed at each replica, example:
  - if  $P_1$  issues the operation  $m_{(1,1)}: x=x+1;$  and
  - If  $P_3$  issues  $m_{(3,1)}: \text{print}(x);$
  - Then, at all replicas  $P_1, P_2, P_3$  execute in the same order, for example:

$\text{print}(x);$

$x=x+1;$

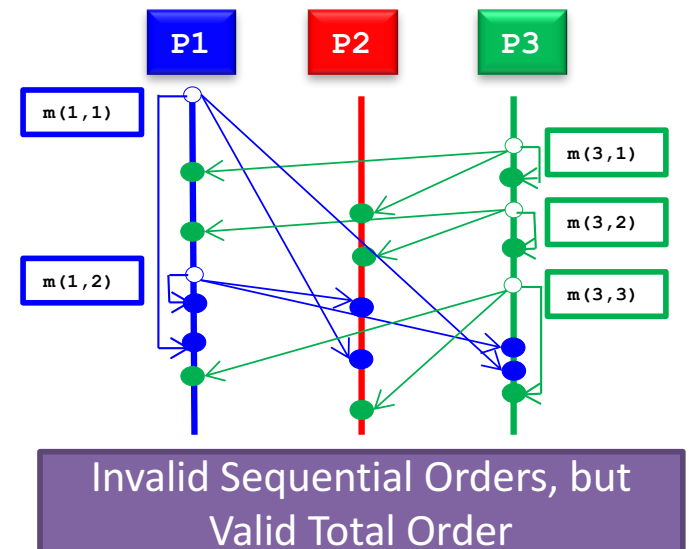



# Types of Ordering

- Total Ordering
- Sequential Ordering
- Causal Ordering

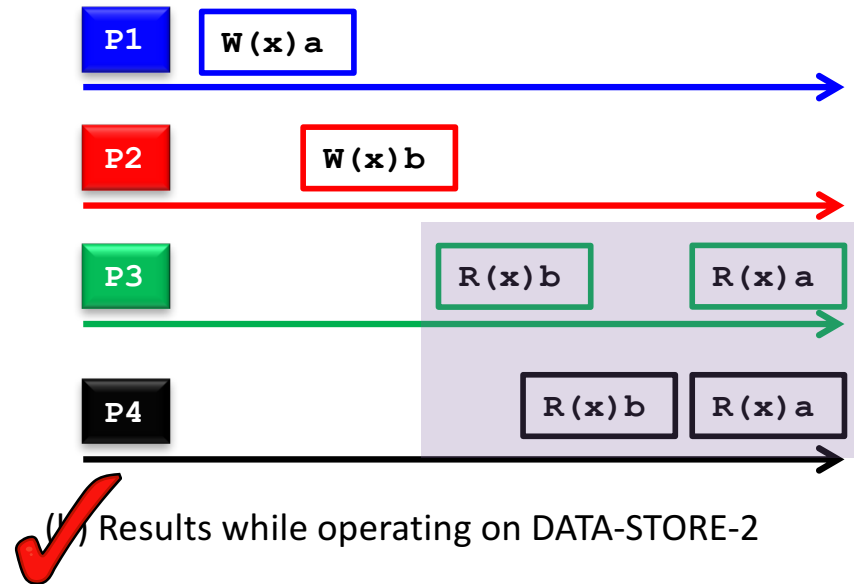
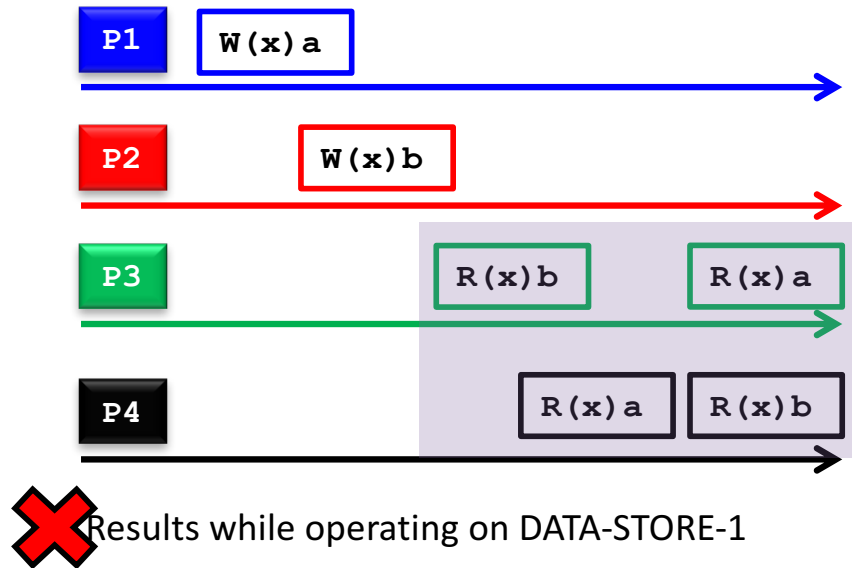
# Sequential Ordering

- + If a process  $P_i$  sends a sequence of messages  $m_{(i,1)}, \dots, m_{(i,n_i)}$ , and
- + Process  $P_j$  sends a sequence of messages  $m_{(j,1)}, \dots, m_{(j,n_j)}$ ,
- + Then, :
  - + At any process, the set of messages received are in the same sequential order
  - + Messages from each individual process appear in this sequence in the order sent by the sender
    - + At every process,  $m_{i,1}$  should be delivered before  $m_{i,2}$ , which is delivered before  $m_{i,3}$  and so on...
    - + At every process,  $m_{j,1}$  should be delivered before  $m_{j,2}$ , which is delivered before  $m_{j,3}$  and so on...



# Sequential Consistency Model


- Sequential Consistency Model enforces that all the update operations are executed at the replicas in a sequential order
  - “all read operations should see writes in the same order”
- Consider a data-store with variable **x** (Initialized to **NULL**)
  - In the two data-stores below, identify the sequentially consistent data-store



**P1** = Process P1       $\rightarrow$  = Timeline at P1       $R(x) b$  = Read variable x; Result is b       $W(x) b$  = Write variable x; Result is b

# Sequential Consistency (cont'd)

- Consider three processes  $P_1$ ,  $P_2$  and  $P_3$  executing multiple instructions on three shared variables  $x$ ,  $y$  and  $z$ 
  - Assume that  $x$ ,  $y$  and  $z$  are set to zero at start




- There are many valid sequences in which operations can be executed at the replica respecting sequential consistency
  - Identify the output?

	<pre>x = 1 print (y,z) y = 1 print (x,z) z = 1 print (x,y)</pre>	<pre>x = 1 y = 1 print (x,z) print (y,z) z = 1 print (x,y)</pre>	<pre>z = 1 print (x,y) print (x,z) y = 1 x = 1 print (y,z)</pre>	<pre>y = 1 z = 1 print (x,y) print (x,z) x = 1 print (y,z)</pre>
Output	001011	101011	000111	010111

# Sequential Cons

<http://olafland.poll daddy.com/s/seqcons>

- Consider three processes  $P_1$ ,  $P_2$  and  $P_3$  and three shared variables  $x$ ,  $y$  and  $z$ 
  - Assume that  $x$ ,  $y$  and  $z$  are set to 1



- There are many valid sequences in the replica respecting sequential consistency
  - Identify the output?
  - Which ones breaks the seq. order?



	<pre>x = 1 print (y,z) y = 1 print (x,z) z = 1 print (x,y)</pre>	<pre>x = 1 y = 1 print (x,z) print (y,z) z = 1 print (x,y)</pre>	<pre>z = 1 print (x,y) print (x,z) y = 1 x = 1 print (y,z)</pre>	<pre>y = 1 z = 1 print (x,y) print (x,z) x = 1 print (y,z)</pre>
Output	001011	101011	000111	010111






<http://olafland.poll daddy.com/s/seqcons>

# Sequential Consistency (cont'd)

- Consider three processes  $P_1$ ,  $P_2$  and  $P_3$  executing multiple instructions on three shared variables  $x$ ,  $y$  and  $z$ 
  - Assume that  $x$ ,  $y$  and  $z$  are set to zero at start



- There are many valid sequences in which operations can be executed at the replica respecting sequential consistency
  - Identify the output?
  - Which ones brakes the seq. order?

	<pre>x = 1 print (y,z) y = 1 print (x,z) z = 1 print (x,y)</pre>	<pre>x = 1 y = 1 print (x,z) print (y,z) z = 1 print (x,y)</pre>	<pre>z = 1 print (x,y) print (x,z) y = 1 x = 1 print (y,z)</pre>	<pre>y = 1 z = 1 print (x,y) print (x,z) x = 1 print (y,z)</pre>
Output	001011	101011	000111	010111

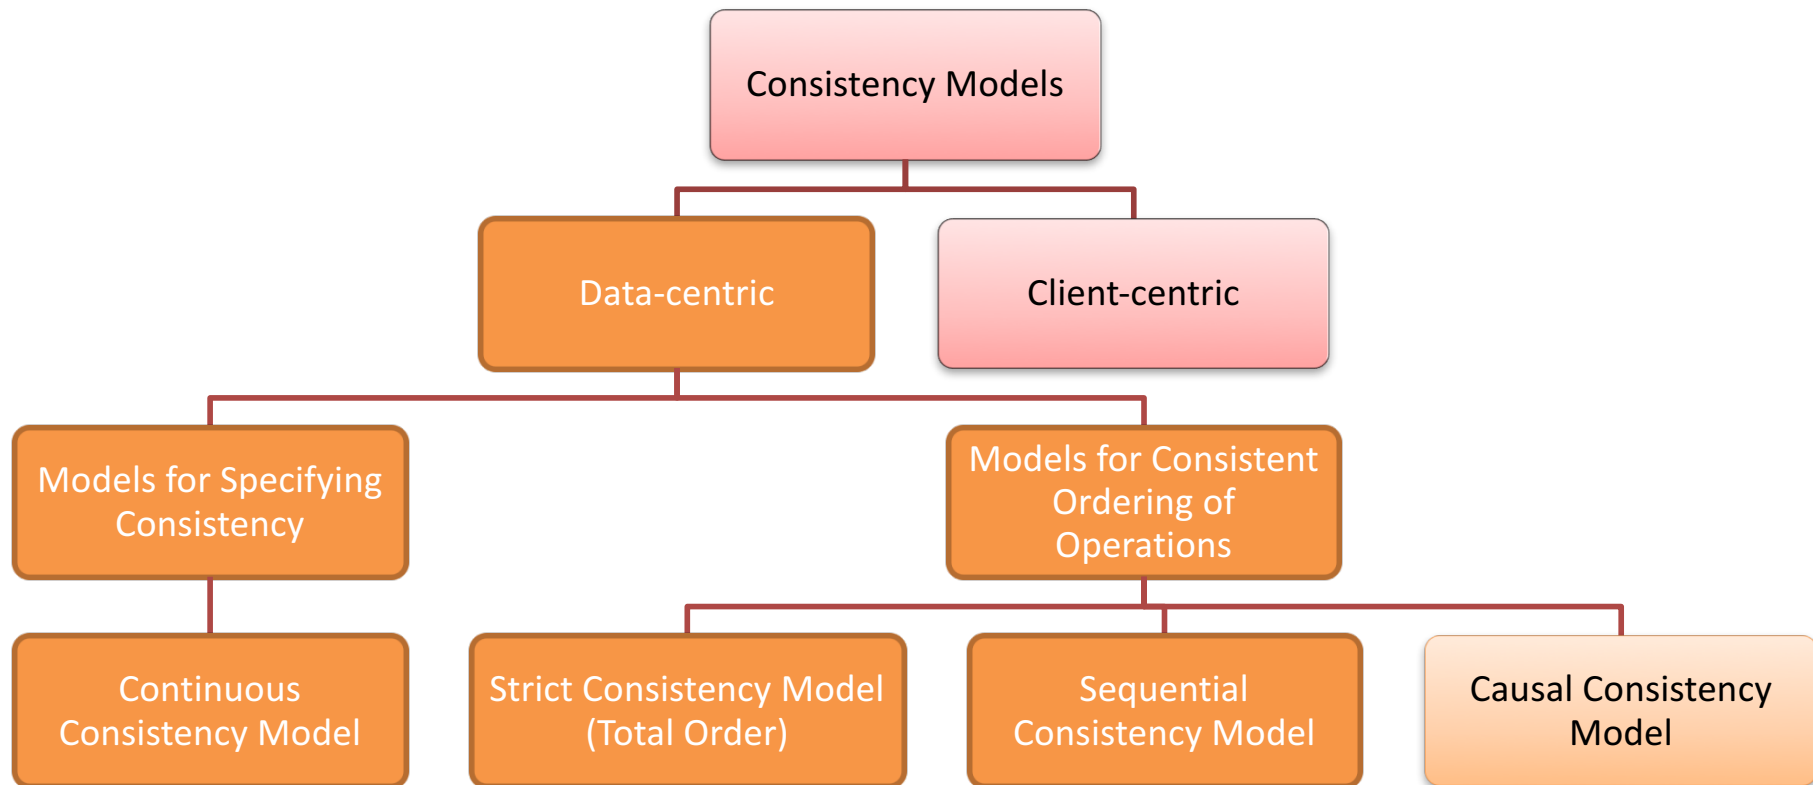


# Implications of Adopting Sequential Consistency Model for Applications

- There might be several different sequentially consistent combinations of ordering
  - Number of combinations for a total of  $n$  instructions =  $O(n!)$
- The contract between the process and the distributed data-store is that the process must accept all of the sequential orderings as valid results
  - A process that works for some of the sequential orderings and does not work correctly for others is INCORRECT

# Recap: Consistency Models

- A consistency model states the level of consistency provided by the *data-store* to the processes while reading and writing the data



# Questions


- Total Order?
  - Same order on all nodes
- Sequential Ordering?
  - At any process, the set of messages received are in the same sequential order
  - Messages from each individual process appear in this sequence in the order sent by the sender

# Causality (Recap)

- Causal relation between two events
  - If **a** and **b** are two events such that **a happened-before b** or  **$a \rightarrow b$** , and
  - If the (logical) time when event **a** and **b** is received at a process  **$P_i$**  is denoted by  **$C_i(a)$**  and  **$C_i(b)$**
  - Then, if we can infer that  **$a \rightarrow b$**  by observing that  **$C_i(a) < C_i(b)$** , then **a** and **b** are causally related
- How can we implement causality?
  - Logical Clocks or Vector Clocks (see prev. lectures)

# Causal vs. Concurrent events


- Consider an interaction between processes  $P_1$  and  $P_2$  operating on replicated data  $x$  and  $y$



Events are causally related

Events are not concurrent

- Computation of  $y$  at  $P_2$  may have depended on value of  $x$  written by  $P_1$   
(As  $P_2$  reads  $x$  before writing  $y$ )



Events are not causally related

Events are concurrent

- Computation of  $y$  at  $P_2$  does not depend on value of  $x$  written by  $P_1$

$P_1$

=Process  $P_1$



=Timeline at  $P_1$

$R(x) b$

=Read variable  $x$ ;  
Result is  $b$

$W(x) b$

= Write variable  $x$ ;  
Result is  $b$

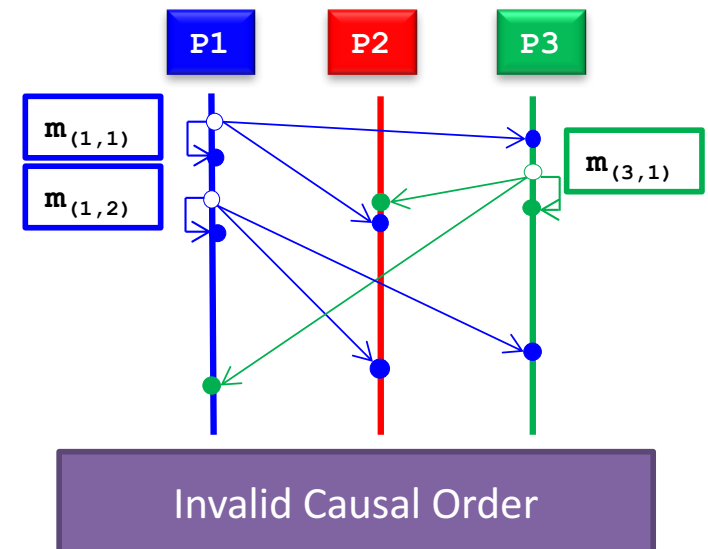

# Causal Ordering

## Causal Order

- If process  $P_i$  sends a message  $m_i$  and  $P_j$  sends  $m_j$ , and if  $m_i \rightarrow m_j$  (operator ' $\rightarrow$ ' is **happened-before** relation) then any correct process that delivers  $m_j$  will deliver  $m_i$  before  $m_j$

Drawback:

- The **happened-before** relation between  $m_i$  and  $m_j$  should be induced before communication

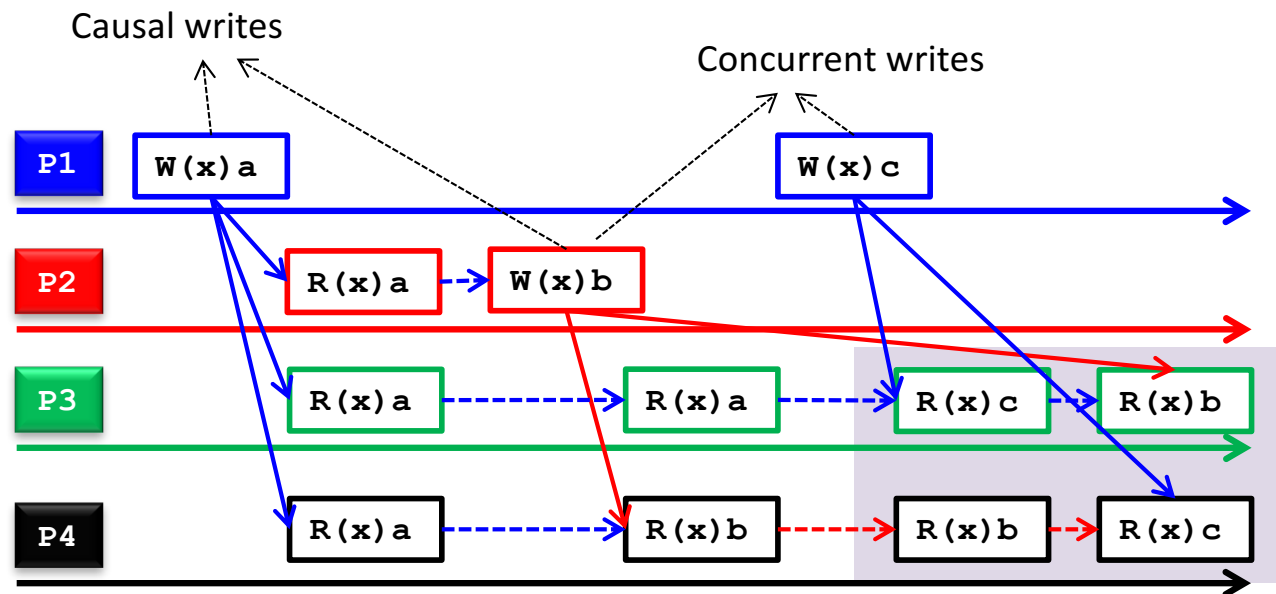




# Causal Consistency Model

- A data-store is causally consistent if:
  - Writes that are potentially causally related must be seen by all the processes in the same order
  - Concurrent writes may be seen in a different order on different machines

# Example of a Causally Consistent Data-store



A Causally Consistent  
Data-Store

But not a Sequentially  
Consistent Data-Store

**P1** = Process P1

$\rightarrow$  = Timeline at P1

**$R(x) b$**  = Read variable  $x$ ;  
Result is  $b$

**$W(x) b$**  = Write variable  $x$ ;  
Result is  $b$

# Causally Consistent Data-Store?


<http://olafland.polldaddy.com/s/causcons>

P1:	W(x)a		W(x)c	
P2:		R(x)a	W(x)b	
P3:		R(x)a		R(x)c
P4:		R(x)a		R(x)b

P1:	W(x)a			
P2:		R(x)a	W(x)b	
P3:			R(x)b	R(x)a
P4:			R(x)a	R(x)b

P1:	W(x)a			
P2:			W(x)b	
P3:			R(x)b	R(x)a
P4:			R(x)a	R(x)b





<http://olafland.polldaddy.com/s/causcons>

# Causally Consistent Data-Store?

P1:	W(x)a		W(x)c	
P2:		R(x)a	W(x)b	
P3:		R(x)a		R(x)c
P4:		R(x)a		R(x)b

This sequence is allowed with a causally-consistent store, but not with sequentially consistent store.

Note:  $W_1(x)a \rightarrow W_2(x)b$ , but not  $W_2(x)b \rightarrow W_1(x)c$

P1:	W(x)a		
P2:		R(x)a	W(x)b
P3:			R(x)b
P4:			R(x)a

A violation of a causally-consistent store.

P1:	W(x)a		
P2:		W(x)b	
P3:			R(x)b
P4:			R(x)a

A correct sequence of events in a causally consistent store.

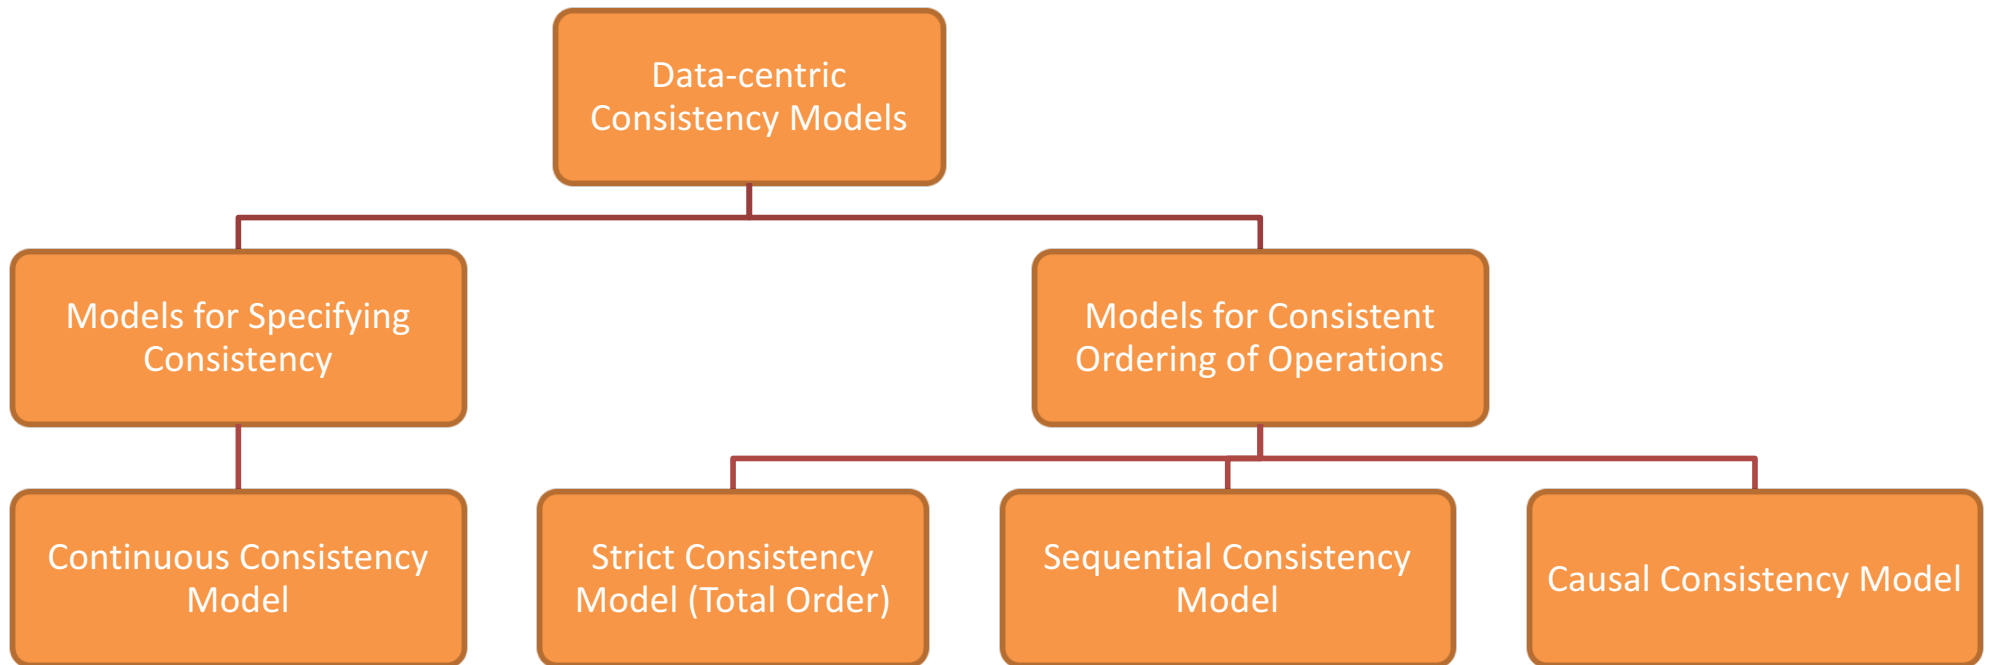
# Implications of adopting a Causally Consistent Data-store for Applications

- Processes have to keep track of which processes have seen which writes
- This requires maintaining a dependency graph between write and read operations
  - Vector clocks provide a way to maintain causal consistency

# Summary

- Total Order?
  - Same order on all nodes
- Sequential Ordering?
  - Same order as sent by the sender
- Causal Order
  - Same order on all nodes only for causally related msg.

# Topics Covered in Data-centric Consistency Models

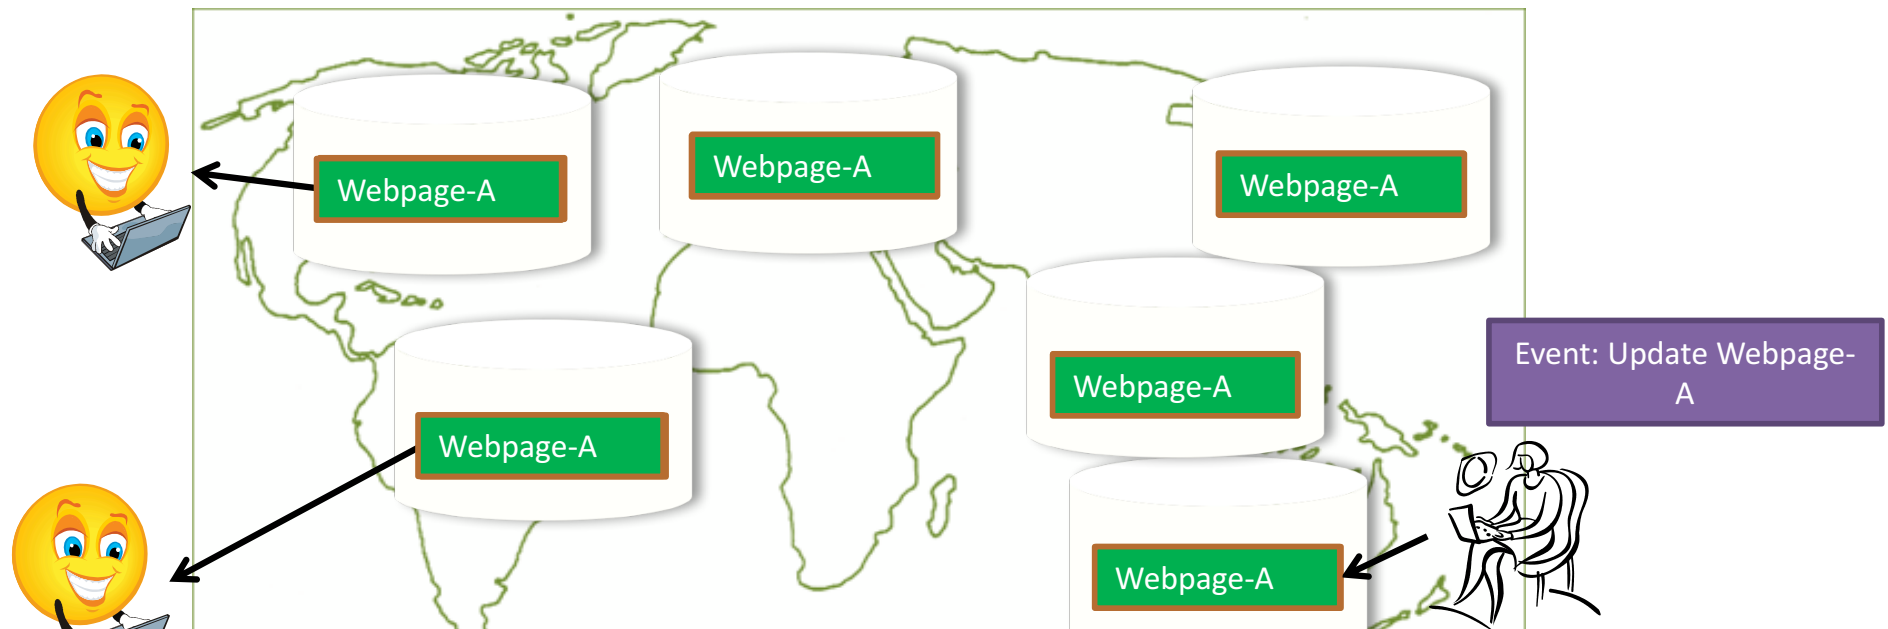


But, is Data-centric Consistency Model good for all applications?



# Applications that can use Data-centric Models

- Data-centric models are applicable when many processes are concurrently updating the data-store
- But, do all applications need all replicas to be consistent?



Data-Centric Consistency Model is too strict when

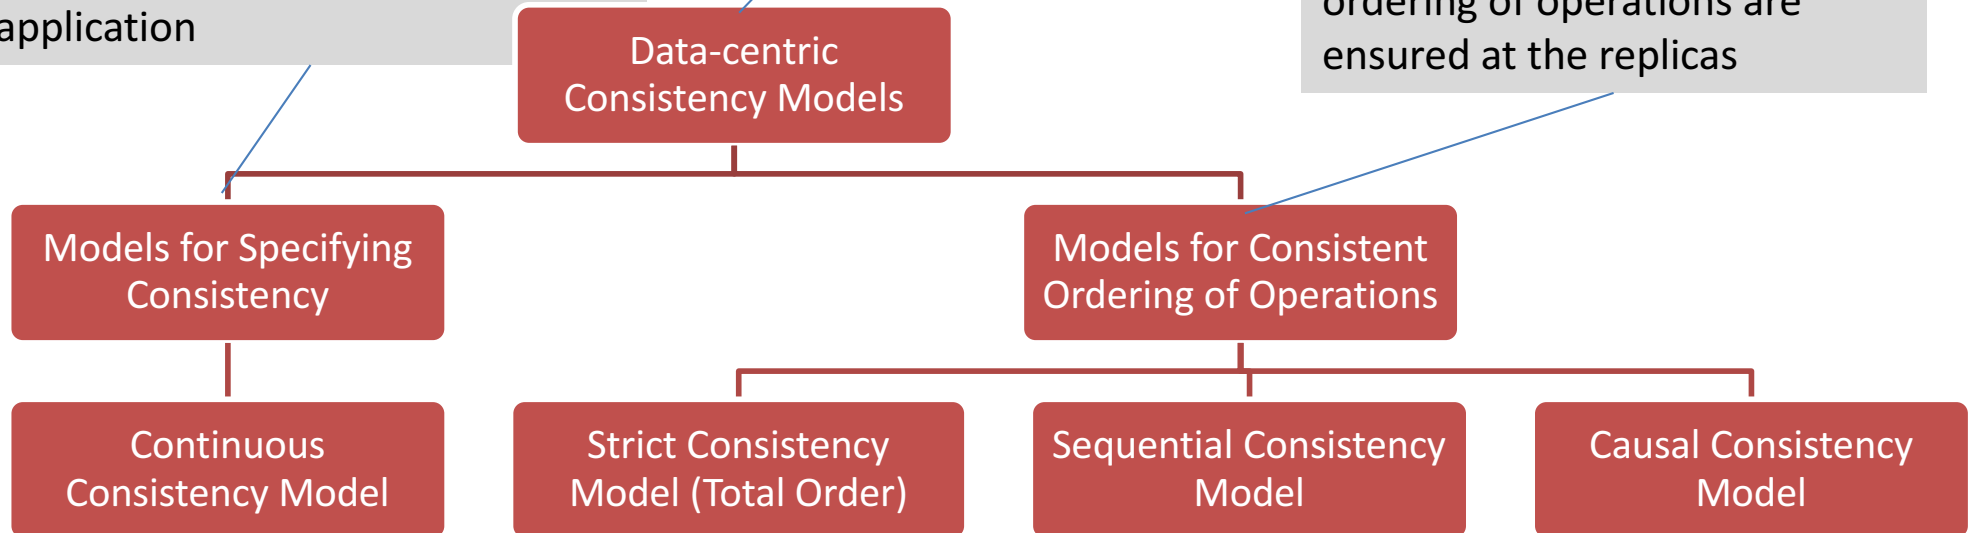
- One client process updates the data
- Other processes read the data, and are OK with reasonably stale data

# Summary of Data-Centric Consistency Models

Data-centric consistency models describe how the replicated data is kept consistent across different data-stores, and what the process can expect from the data-store

These models allow measuring and specifying the consistency levels that are tolerable to the application

These models specify what ordering of operations are ensured at the replicas



Data-centric models are too strict when:

- most operations are read operations
- updates are generally triggered from one client process

# Summary

- Replication is necessary for improving performance, scalability and availability, and for providing fault-tolerance
- Replicated data-stores should be designed after carefully evaluating the trade-off between tolerable data inconsistency and efficiency
- Consistency Models describe the contract between the data-store and process about what form of consistency to expect from the system
- Data-centric consistency models:
  - Continuous Consistency Models provide mechanisms to measure and specify inconsistencies
  - Consistency Models can be defined based on the type of ordering of operations that the replica guarantees the application

# Next Classes

- Consistency Models
  - Client-Centric Consistency Models
- Replica Management
  - Replica management studies:
    - when, where and by whom replicas should be placed
    - which consistency model to use for keeping replicas consistent
- Consistency Protocols
  - We study selected implementations of consistency models

# Questions?

In part, inspired from / based on slides from

- Vinay Kolar
- A. S. Tanenbaum, M. v. Steen